

# Operating Systems and Networks

Lecture 03:

Introduction to OS-part 2

Behzad Bordbar

# Recap

## Lecture 1:

- ❑ CPU: how it works, user and kernel mode, use of register, cache, ...
- ❑ System call trap and interrupt
- ❑ What happens when computer starts?
- ❑ Different types of Memory and their speed

Will continue with preliminaries of OS

# Demo lecture

## Contents

- ❑ Service view (provider of services) of the OS
- ❑ Shell
- ❑ Everything a directory
- ❑ mkdir, mv, cp,...
- ❑ Access control
- ❑ Find, grep
- ❑ |, >, >>, ; and their differences
- ❑ wget,...

# Contents

- ❑ How does mouse and keyboard work?
- ❑ Device controller
- ❑ CPU multitasking
- ❑ Time sharing
- ❑ a short study of system calls
  - ❑ API

# How does mouse, keyboard ...work?

- ❑ We said

Hardware may trigger an interrupt at any time by sending a signal to the CPU.

- ❑ But how?

- ❑ **Devices** interact via **device controller** connected through a common **bus** to CPU. **Draw picture!**

- ❑ small computer-systems interface (**SCSI**) controller.

- ❑ SCSI controller: hardware (card or chip) that allows SCSI storage device to communicate with the operating system using a SCSI bus

# Device controller

- ❑ Maintains some local buffer storage and a set of special-purpose registers.
- ❑ Device controller moves the data between the peripheral devices that it controls and its local buffer storage.
- ❑ Operating systems have a device driver for each device controller.
- ❑ **you download drivers!** Manually or automatically
- ❑ Device driver understands the device controller and provides the rest of the OS with a uniform interface to the device: copy the same no matter what device

## Back to, How does I/O (mouse..) work?

- ❑ Device driver loads the appropriate registers within the device controller.
- ❑ Device controller examines the contents of the registers to determine what action to take (read char from k/b)
- ❑ controller starts the transfer of data from the device to its local buffer.
- ❑ when transfer of data is complete, the device controller informs the device driver via an interrupt that it has finished its operation.
- ❑ Device driver then returns control to the operating system, possibly returning the data or a pointer to the data if the operation was a read or status (success...)
- ❑ Exercise: draw a sequence diagram for yourself!

## How does I/O (mouse..) work? (cont..)

- ❑ This form of interrupt-driven I/O is fine for moving **small amounts** of data
- ❑ Not suitable for bulk data movement such as disk I/O.  
Instead: Direct Memory Access (DMA) is used that takes CPU out of the loop.
- ❑ After setting up buffers, pointers, and counters for the I/O device, the device controller transfers an entire block of data directly to or from its own buffer storage to memory, with no intervention by the CPU.
- ❑ Hence: only one interrupt is generated per block, to tell the device driver that the operation has completed, rather than the one interrupt per byte generated for low-speed devices.
- ❑ CPU is made free to do other things



# Multitasking in CPU

- ❑ OS picks and begins to execute one of the jobs in memory.
- ❑ Eventually job may have to wait for some task, such as an I/O operation, to complete.
- ❑ OS switches to, and executes, another job.
- ❑ When that job needs to wait, the CPU switches to another job, and so on.
- ❑ Eventually, the first job finishes waiting and gets the CPU back.

## Time-sharing:

- ❑ CPU executes multiple jobs by switching among them, but the switches occur so frequently that the users can interact with each program while it is running.

# Time sharing

requires CPU **scheduling** of user tasks. But how?

- ❑ Each user has at least one separate program in memory.
- ❑ A program loaded into memory and executing is called a **process**. We will study this in details!
- ❑ When a process executes, it typically executes for only a short time before it either finishes or needs to perform I/O.
- ❑ I/O takes long long long time compare to execution! (look at the speed of access slides!)

# Time sharing

- ❑ Time sharing: several jobs be kept simultaneously in memory.
- ❑ CPU scheduling: process of deciding which job is brought to memory to be executed, when there are not enough room.

Reasonable response time must be ensured:

1. processes are **swapped** in and out of main memory to the disk
  2. use **virtual memory**: a technique that allows the execution of a process that is not completely in memory
- ❑ virtual-memory scheme enables users to run programs that are larger than actual **physical memory**. Further, it abstracts main memory into a large, uniform array of storage, separating **logical memory** as viewed by the user from physical memory. This arrangement frees programmers from concern over memory-storage limitations.

# Dual mode




Figure from Dragon book.

- ❑ Dual mode OS protect from harm caused by privileged instructions
- ❑ Extended to multi mode by domains: Dom0, DomU

# Why do we need hardware support?

- ❑ MS-DOS: Intel 8088 architecture, which has no mode bit
- ❑ user program can wipe out the whole OS
- ❑ programs are able to write to a device ...

In dual mode:

- ❑ hardware detects errors that violate modes and handle them by OS
- ❑ stops user program attempts to execute an illegal instruction or to access memory of other users

When error detected

- ❑ OS must terminate the program
- ❑ OS gives error message
- ❑ produces memory dumps by writing to a file (users can check or OS vendors can check (Sun)).

# system calls

provide an interface to the services made available by an operating system.

What language: are System-call written in?

- ❑ typically C and C++ and sometimes assembly-language involved

- ❑ explain the system calls for reading data from one file and writing to another file:

`$cp file1 file2`

open file1, possible error(print, abort), create file2 (file2 exists, rewrite/rename...), start read and write (errors:disk space, memory stick unplugged...), all read and written, close files, ack

`Do I access system call directly?`

# system calls: API to wrap system calls

## Application Programming Interface (API)

- ❑ specifies a set of functions that are available to an application programmer, including the **parameters** that are passed to each function and the **return values** the programmer can expect.
- ❑ programmer accesses an API via a library of code provided by the operating system.

## Example of APIs:

### 1. Windows API for Windows systems

Example: `CreateProcess()` which invokes the `NTCreateProcess()` system call in the Windows kernel  
return value 0 or 1 (error)

# system calls: API (continue)

Example of APIs:

2. POSIX API for POSIX-based systems (UNIX, Linux, and Mac OS X)

- ❑ programmer accesses an API via a library of code provided by the operating system.

Example: read

input:

- ❑ int fd: file descriptor to be read
- ❑ void \*buf: pointer into buffer to be read into
- ❑ size\_t count: maximum number of bytes to read

output:

- ❑ number of bytes read (if success)
- ❑ -1 if fail
- ❑ UNIX and Linux for programs written in the C language, the library is called libc.



## system calls: API (continue)

Example of APIs:

3. Java API for programs that run on the Java virtual machine.

`getParentFile()`

invoked on a file object.

output:

Returns the abstract pathname of this abstract pathname's parent, or null if this pathname does not name a parent directory.

JVM uses the OS system calls.

# why do we use API?

Why not invoking actual system calls directly?

- ❑ Program portability: program can compile and run on any system that supports the API
- ❑ system calls can often be more detailed and difficult to work with
- ❑ give access to high level objects (java API)

Do you know interfaces in java? using system calls is like implementing an (or many) interfaces

## What happens when a user prog. makes a system call

- ❑ caller only needs to know the signature!
- ❑ method call and parameters are passed into a registers
- ❑ values saved in memory for example on table or stacks but addresses in registers
- ❑ Stack is preferred because do not put limit on the number of parameters stored.

## summary

- ❑ Study of I/O devices
- ❑ Device Controller
- ❑ How OS manages multiple tasks
- ❑ System call and their use in user programs
- ❑ API and examples of API
- ❑ Similarity between API and interface